

Untitled5

April 7, 2021

```
[1]: #week 1 task

#Importing all necessary libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

UsageError: Line magic function `%matplotlib` not found.

```
[ ]:
```

```
[2]: data = pd.read_csv('health care diabetes.csv') #importing CSV file
```

```
[3]: data.head()
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: data.isnull().any() #checking for Null data
```

```
[4]: Pregnancies           False
      Glucose              False
      BloodPressure        False
      SkinThickness        False
```

```

Insulin                False
BMI                    False
DiabetesPedigreeFunction  False
Age                    False
Outcome                False
dtype: bool

```

```
[5]: data.info() #data Information
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[6]: Positive = data[data['Outcome']==1] #grouping the data where outcome is 1
Positive.head(5)
```

```
[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
2	8	183	64	0	0	23.3	
4	0	137	40	35	168	43.1	
6	3	78	50	32	88	31.0	
8	2	197	70	45	543	30.5	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
2	0.672	32	1
4	2.288	33	1
6	0.248	26	1
8	0.158	53	1


```
[7]: data['Glucose'].value_counts().head(7)
```

```
[7]: 100    17
      99    17
```

```
129    14
125    14
111    14
106    14
95     13
Name: Glucose, dtype: int64
```

```
[8]: plt.hist(data['Glucose'])
```

```
[8]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
      array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
            179.1, 199. ]),
      <BarContainer object of 10 artists>)
```




```
[9]: data['BloodPressure'].value_counts().head(7)
```

```
[9]: 70    57
      74    52
      68    45
      78    45
      72    44
      64    43
      80    40
      Name: BloodPressure, dtype: int64
```

```
[10]: plt.hist(data['BloodPressure'])
```

```
[10]: (array([ 35.,  1.,  2., 13., 107., 261., 243.,  87., 14.,  5.]),
      array([ 0. , 12.2, 24.4, 36.6, 48.8, 61. , 73.2, 85.4, 97.6,
            109.8, 122. ]),
      <BarContainer object of 10 artists>)
```




```
[11]: data['SkinThickness'].value_counts().head(7)
```

```
[11]: 0      227
      32      31
      30      27
      27      23
      23      22
      33      20
      18      20
      Name: SkinThickness, dtype: int64
```

```
[12]: plt.hist(data['SkinThickness'])
```

```
[12]: (array([231., 107., 165., 175.,  78.,  9.,  2.,  0.,  0.,  1.]),
      array([ 0. ,  9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),
      <BarContainer object of 10 artists>)
```




```
[13]: data['Insulin'].value_counts().head(7)
```

```
[13]: 0      374
      105     11
      140      9
      130      9
      120      8
      100      7
       94      7
      Name: Insulin, dtype: int64
```

```
[14]: plt.hist(data['Insulin'])
```

```
[14]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),
      array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,
              761.4, 846. ]),
      <BarContainer object of 10 artists>)
```




```
[15]: data['BMI'].value_counts().head(7)
```

```
[15]: 32.0    13
      31.6    12
      31.2    12
      0.0    11
      33.3    10
      32.4    10
      32.8     9
      Name: BMI, dtype: int64
```

```
[16]: plt.hist(data['BMI'])
```

```
[16]: (array([ 11.,  0., 15., 156., 268., 224., 78., 12.,  3.,  1.]),
      array([ 0. ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
              60.39, 67.1 ]),
      <BarContainer object of 10 artists>)
```



```
[9]: data.describe().transpose()
```


```
[9]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	
Insulin	768.0	79.799479	115.244002	0.000	0.00000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00


```
[10]: plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)
```

```
[10]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
            14.,  7.,  3.,  1.,  1.,  0.,  1.]),
       array([ 0.   ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
            26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
            53.68 , 57.035, 60.39 , 63.745, 67.1  ]),
       [<matplotlib.patches.Polygon at 0x7f7aa9bcd50>])
```



```
[11]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

```
[11]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
            27., 29., 22., 24., 21., 25., 14.]),
       array([ 0.   ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
            79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
            159.2 , 169.15, 179.1 , 189.05, 199.  ]),
       [<matplotlib.patches.Polygon at 0x7f7aa9aea1d0>])
```





```
[12]: Positive['Glucose'].value_counts().head(7)
```

```
[12]: 125      7
      158      6
      128      6
      115      6
      129      6
      146      5
      162      5
      Name: Glucose, dtype: int64
```

```
[13]: plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
[13]: (array([16.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  6.,  6., 19., 37., 56.,
          36., 41., 31.,  7.,  4.,  4.,  3.]),
      array([ 0. ,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
          51.3, 57. , 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
          102.6, 108.3, 114. ]),
      [<matplotlib.patches.Polygon at 0x7f7aa9a5f510>])
```




```
[14]: Positive['BloodPressure'].value_counts().head(7)
```

```
[14]: 70    23
      76    18
      78    17
      74    17
      72    16
      0    16
      82    13
      Name: BloodPressure, dtype: int64
```

```
[15]: plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
[15]: (array([88.,  1.,  4., 10., 18., 30., 41., 34., 23., 15.,  1.,  1.,  1.,
          0.,  0.,  0.,  0.,  0.,  0.,  1.]),
      array([ 0. ,  4.95,  9.9 , 14.85, 19.8 , 24.75, 29.7 , 34.65, 39.6 ,
          44.55, 49.5 , 54.45, 59.4 , 64.35, 69.3 , 74.25, 79.2 , 84.15,
          89.1 , 94.05, 99.  ]),
      [<matplotlib.patches.Polygon at 0x7f7aa9a477d0>])
```




```
[16]: Positive['SkinThickness'].value_counts().head(7)
```

```
[16]: 0      88
      32     14
      33      9
      30      9
      39      8
      35      8
      36      8
      Name: SkinThickness, dtype: int64
```

```
[17]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```

```
[17]: (array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,
          5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
      array([ 0. , 42.3, 84.6, 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,
          380.7, 423. , 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,
          761.4, 803.7, 846. ]),
      [<matplotlib.patches.Polygon at 0x7f7aa99b1410>])
```




```
[18]: Positive['Insulin'].value_counts().head(7)
```


```
[18]: 0      138
      130      6
      180      4
      156      3
      175      3
      194      2
      125      2
      Name: Insulin, dtype: int64
```

```
[19]: BloodPressure = Positive['BloodPressure']
      Glucose = Positive['Glucose']
      SkinThickness = Positive['SkinThickness']
      Insulin = Positive['Insulin']
      BMI = Positive['BMI']
```


```
[20]: plt.scatter(BloodPressure, Glucose, color=['b'])
      plt.xlabel('BloodPressure')
      plt.ylabel('Glucose')
      plt.title('BloodPressure & Glucose')
      plt.show()
```




```
[21]: g =sns.scatterplot(x= "Glucose" ,y= "BloodPressure",  
                        hue="Outcome",  
                        data=data);
```



```
[22]: B =sns.scatterplot(x= "BMI" ,y= "Insulin",  
                        hue="Outcome",  
                        data=data);
```



```
[23]: S =sns.scatterplot(x= "SkinThickness" ,y= "Insulin",  
                        hue="Outcome",  
                        data=data);
```



```
[24]: data.corr()
```

```
[24]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	


	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	

```
Age Outcome
```

Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000


```
[25]: sns.heatmap(data.corr())
```

```
[25]: <AxesSubplot:>
```



```
[26]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True,cmap='viridis')
```

```
[26]: <AxesSubplot:>
```

```
[27]: ##logistic regresion
```

```
[28]: data.head(5)
```

```
[28]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6     148           72           35         0  33.6
1           1      85           66           29         0  26.6
2           8     183           64            0         0  23.3
3           1      89           66           23        94  28.1
4           0     137           40           35       168  43.1

   DiabetesPedigreeFunction  Age  Outcome
```

0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[31]: features = data.iloc[:,[0,1,2,3,4,5,6,7]].values
      label = data.iloc[:,8].values
```

```
[32]: #Train test split
      from sklearn.model_selection import train_test_split
      X_train,X_test,y_train,y_test = train_test_split(features,
                                                         label,
                                                         test_size=0.3,
                                                         random_state =10)
```

```
[33]: #Create model
      from sklearn.linear_model import LogisticRegression
      model = LogisticRegression()
      model.fit(X_train,y_train)
```

/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:940:
 ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
[33]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                        intercept_scaling=1, l1_ratio=None, max_iter=100,
                        multi_class='auto', n_jobs=None, penalty='l2',
                        random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                        warm_start=False)
```

```
[34]: print(model.score(X_train,y_train))
      print(model.score(X_test,y_test))
```

0.7821229050279329
 0.7402597402597403

```
[35]: from sklearn.metrics import confusion_matrix
      cm = confusion_matrix(label,model.predict(features))
      cm
```

```
[35]: array([[443,  57],
           [120, 148]])
```

```
[36]: from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```


	precision	recall	f1-score	support
0	0.79	0.89	0.83	500
1	0.72	0.55	0.63	268
accuracy			0.77	768
macro avg	0.75	0.72	0.73	768
weighted avg	0.76	0.77	0.76	768

```
[37]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.836

```
[37]: [<matplotlib.lines.Line2D at 0x7f7aa405b150>]
```



```
[38]: #Applying Decision Tree Classifier
from sklearn.tree import DecisionTreeClassifier
model3 = DecisionTreeClassifier(max_depth=5)
model3.fit(X_train,y_train)
```

```
[38]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                             max_depth=5, max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=None, splitter='best')
```

```
[39]: model3.score(X_train,y_train)
```

```
[39]: 0.845437616387337
```

```
[40]: model3.score(X_test,y_test)
```

```
[40]: 0.7316017316017316
```

```
[41]: #Applying Random Forest
from sklearn.ensemble import RandomForestClassifier
model4 = RandomForestClassifier(n_estimators=11)
model4.fit(X_train,y_train)
```

```
[41]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=None, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=11,
                             n_jobs=None, oob_score=False, random_state=None,
                             verbose=0, warm_start=False)
```

```
[42]: model4.score(X_train,y_train)
```

```
[42]: 0.9832402234636871
```

```
[43]: model4.score(X_test,y_test)
```

```
[43]: 0.7142857142857143
```

```
[44]: #Support Vector Classifier

from sklearn.svm import SVC
model5 = SVC(kernel='rbf',
              gamma='auto')
model5.fit(X_train,y_train)
```

```
[44]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
         decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
         max_iter=-1, probability=False, random_state=None, shrinking=True,
         tol=0.001, verbose=False)
```

```
[45]: model5.score(X_train,y_train)
```

```
[45]: 1.0
```

```
[46]: model5.score(X_test,y_test)
```

```
[46]: 0.6233766233766234
```

```
[47]: #Applying K-NN
from sklearn.neighbors import KNeighborsClassifier
model2 = KNeighborsClassifier(n_neighbors=7,
                             metric='minkowski',
                             p = 2)
model2.fit(X_train,y_train)
```

```
[47]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                          metric_params=None, n_jobs=None, n_neighbors=7, p=2,
                          weights='uniform')
```


```
[48]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".
      ↪format(tpr,fpr,thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
```

AUC: 0.832

```
True Positive Rate - [0.          0.05597015 0.23507463 0.44776119 0.60447761
0.73880597
0.87313433 0.98134328 1.          ], False Positive Rate - [0.          0.          0.026
0.058 0.128 0.252 0.406 0.65   1.          ] Thresholds - [2.          1.
0.85714286 0.71428571 0.57142857 0.42857143
0.28571429 0.14285714 0.          ]
```

```
[48]: Text(0, 0.5, 'True Positive Rate')
```




```
[49]: #Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.626 auc=0.727 ap=0.728

[49]: [



```
[50]: #Precision Recall Curve for KNN


from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
```



```
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.'))
```

f1=0.656 auc=0.740 ap=0.697

[50]: [<matplotlib.lines.Line2D at 0x7f7a9ad4dc10>]



[51]: *#Precision Recall Curve for Decission Tree Classifier*

```
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model3.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model3.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
```


```

# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.684 auc=0.790 ap=0.764

[51]: [<matplotlib.lines.Line2D at 0x7f7a9ac5e550>]



```

[52]: #Precision Recall Curve for Random Forest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model4.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)

```


```

# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.853 auc=0.929 ap=0.916

[52]: [<matplotlib.lines.Line2D at 0x7f7a9abc6a50>]



[]: